

Using Two-Sided Blocking Point to Point Communications in MPI

Erik Saule

`esaule@uncc.edu`

Parallel and Distributed Computing

Learning Outcomes

After this lecture you will be able to

- Use blocking point to point communication primitives.

Outline

- 1 Idea
- 2 Syntax
- 3 Underspecifying MPI_Recv
- 4 Further

Two-sided communications

What?

- Make a communication between a pair of processes
- The sender makes a call to send
- The receiver makes a different call to receive
- This breaks process symmetry :(

Why?

- Most engineering simulation
- Manager Worker Systems like Folding at Home
- Any application structure with dedicated nodes (database, IO node, specialized units like GPUs)

Outline

- 1 Idea
- 2 Syntax
- 3 Underspecifying MPI_Recv
- 4 Further

MPI_Send

man MPI_Send

MPI_Send

man MPI_Recv

Simple example

```
#include <mpi.h>
#include <iostream>
#include <vector>

void printvector(const std::string & msg, const std::vector<int>& v) {
    std::cout<<msg;
    for (auto& a : v)
        std::cout<<a<<" ";
    std::cout<<"\n";
}

int main (int argc, char*argv[]) {
    MPI_Init (&argc, &argv);
    int size;
    int rank;
    MPI_Comm_size(MPI_COMM_WORLD, &size);
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);

    std::cout<<"I am "<<rank<<" out of "<<size<<"\n";

    int vecsize = 3;

    std::vector<int> vec(vecsize);
    std::vector<int> vec2(vecsize);

    for (int i=0; i<vecsize; ++i)
        vec[i] = rank*100+i;

    printvector("before: ", vec);

    if (rank == 0) {
        MPI_Send(&(vec[0]), vecsize, MPI_INT, 1, 100, MPI_COMM_WORLD);
    }
```


Outline

- 1 Idea
- 2 Syntax
- 3 Underspecifying MPI_Recv**
- 4 Further

What's up with MPI_Status?

Because recv takes wildcards

- Source can be MPI_ANY_SOURCE
- Tag can be MPI_ANY_TAG
- Count just needs to be larger than what is being sent

How to know the unknown?

- status.MPI_SOURCE
- status.MPI_TAG
- `int MPI_Get_count(const MPI_Status *status, MPI_Datatype datatype, int *count)`

Usage example of status

```
#include <mpi.h>
#include <iostream>
#include <vector>

void printvector(const std::string & msg, const std::vector<int>& v) {
    std::cout<<msg;
    for (auto& a : v)
        std::cout<<a<<" ";
    std::cout<<"\n";
}

int main (int argc, char*argv[]) {
    MPI_Init (&argc, &argv);

    int size;
    int rank;

    MPI_Comm_size(MPI_COMM_WORLD, &size);
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);

    std::cout<<"I am "<<rank<<" out of "<<size<<"\n";

    int vecsize = 3;

    std::vector<int> vec(vecsize);
    std::vector<int> vec2(vecsize*10);

    for (int i=0; i<vecsize; ++i)
        vec[i] = rank*100+i;

    printvector("before: ", vec);
```

Outline

- 1 Idea
- 2 Syntax
- 3 Underspecifying MPI_Recv
- 4 Further

External

Books:

- Using MPI, 3rd edition. William Gropp, Ewing Lusk and Anthony Skjellum. MIT Press. Available through the library at <https://librarylink.uncc.edu/login?url=http://ieeexplore.ieee.org/xpl/bkabstractplus.jsp?bkn=6981847>
- Using Advanced MPI. William Gropp, Torsten Hoefler, Rajeev Thakur and Ewing Lusk. MIT Press. Available through the library at <https://librarylink.uncc.edu/login?url=http://ieeexplore.ieee.org/xpl/bkabstractplus.jsp?bkn=6981848>

MPI implementations:

- MPICH <https://www.mpich.org>
- OpenMPI <https://www.open-mpi.org/>

API Documentation:

- MPICH man pages <https://www.mpich.org/static/docs/v3.2/www3/index.htm>
- OpenMPI documentation <https://www.open-mpi.org/doc/v3.0/>

Slides from colleagues:

- Tutorial on MPI programming. Victor Eijkhout.
<https://bitbucket.org/VictorEijkhout/parallel-computing-book/raw/e11748c8d8ae874ed645566ba0e82aa787ecf959/EijkhoutMPIlecture.pdf>
- MPI for Dummies. Pavan Balaki, Torsten Hoefler.
https://htor.inf.ethz.ch/teaching/mpi_tutorials/ppopp13/2013-02-24-ppopp-mpi-basic.pdf

Tutorial:

- part one of Parallel Programming in MPI and OpenMP. Victor Eijkhout. Draft at <https://bitbucket.org/VictorEijkhout/parallel-computing-book/raw/e11748c8d8ae874ed645566ba0e82aa787ecf959/EijkhoutParComp.pdf>